



Performance Evaluation of WebAssembly

Punnal Khan and Md Shafiuzzaman

- “Assembly” like language for the web
- Compilation target for other languages
- Delivered as compiled binaries
- Execution environment: browser (Web APIs) or a virtual machine
- Memory model: loads and stores to untyped arrays of bytes
- Modular secure compilation

C++

```
int add(){  
    int b = 9;  
    int a = 1 + b;  
  
    return a;  
}
```

WebAssembly Text

```
(module  
  (memory 1)  
  (func $add (result i32)  
    (local $var0 i32)  
    i32.const 0  
    i32.load offset=4  
    i32.const 16  
    i32.sub  
    tee_local $var0  
    :  
    :  
    :  
  )  
)
```

WebAssembly Bytes

```
00000000: 0061 736d 0100  
0000 0185 8080 8000 0160  
00000010: 0001 7f03 8380  
8080 0002 0000 0484 8080  
00000020: 8000 0170 0000  
0583 8080 8000 0100 0106  
00000030: 8180 8080 0000  
079b 8080 8000 0306 6d65  
00000040: 6d6f 7279 0200  
:  
:  
:
```

- Node.js C++ Addon: update requires for every new version of Node.js
- Asm.js: still a javascript

WebAssembly claims..

- Runs faster than JavaScript
- Near-native performance

Group, WebAssembly Community, and Andreas Rossberg. “WebAssembly Specification,” n.d., 185.

- A barcode scanner implemented in WebAssembly at eBay boosted the performance of the JavaScript implementation by 50 times [1]
- WebAssembly is slower than JavaScript on the Samsung Internet browser (v7.2.10.12) when performing multiplications on matrices of certain sizes [2]

[1] Senthil, Padmanabhan and Pranav Jha. 2020. WebAssembly at eBay: A Real-World Use Case.
<https://tech.ebayinc.com/engineering/webassembly-at-ebay-a-real-world-use-case/>

[2] Winston Chen. 2018. Performance Testing Web Assembly vs JavaScript.
<https://medium.com/samsung-internet-dev/performance-testing-web-assembly-vs-javascript-e07506fd5875>

Is WebAssembly fast?

```
hira@hira-Inspiron-3443:~/WebAssembly/iswasmfast$ node benchmark.js
Levenstein Distance:
  Native x 212,945 ops/sec  $\pm 1.25\%$  (88 runs sampled)
  Web Assembly x 190,769 ops/sec  $\pm 0.47\%$  (92 runs sampled)
  Fastest is Native

Fibonacci:
  Native x 4,909,081 ops/sec  $\pm 0.45\%$  (94 runs sampled)
  Web Assembly x 13,347,474 ops/sec  $\pm 2.62\%$  (89 runs sampled)
  Fastest is Web Assembly

Fermat Primality Test:
  Native x 2,149,635 ops/sec  $\pm 0.71\%$  (90 runs sampled)
  Web Assembly x 3,083,462 ops/sec  $\pm 0.96\%$  (86 runs sampled)
  Fastest is Web Assembly

Simple Linear Regression:
  Native x 253,335 ops/sec  $\pm 1.25\%$  (91 runs sampled)
  Web Assembly x 38,706 ops/sec  $\pm 0.71\%$  (89 runs sampled)
  Web Assembly using TypedArrays x 47,491 ops/sec  $\pm 2.16\%$  (88 runs sampled)
  Fastest is Native

SHA256:
  Native x 15,145 ops/sec  $\pm 21.00\%$  (90 runs sampled)
  Web Assembly x 42,139 ops/sec  $\pm 0.76\%$  (90 runs sampled)
  Fastest is Web Assembly
```

Courtesy: <https://github.com/zandaqo/iswasmfast>

Benchmarks: 41 WebAssembly binaries and 41 JavaScript programs compiled from 41 widely-used C benchmarks (PolyBenchC)

| Input Size | XS | S | M | L | XL |
|--------------------------|-------------------------|-------------------------|----------------------------|-------------------------------|-------------------------------|
| Input size in benchmarks | MINI_DATASET | SMALL_DATASET | MEDIUM_DATASET | LARGE_DATASET | EXTRALARGE_DATASET |
| CORRELATION | M=28 N =32 | M=80 N=100 | M=240 N=260 | M=1200 N=1400 | M=2600 N=3000 |
| COVARIANCE | M=28 N =32 | M=80 N=100 | M=240 N=260 | M=1200 N=1400 | M=2600 N=3000 |
| GEMM | NI=20 NJ=25 NK=30 | NI=60 NJ=70 NK=80 | NI=200 NJ=220 NK=240 | NI=1000 NJ=1100 NK=1200 | NI=2000 NJ=2300 NK=2600 |
| GEMVER | N=40 | N=120 | N=400 | N=2000 | N=4000 |

Yan, Yutian, Tengfei Tu, Lijian Zhao, Yuchen Zhou, and Weihang Wang. “Understanding the Performance of Webassembly Applications.” In Proceedings of the 21st ACM Internet Measurement Conference, 533–49. Virtual Event: ACM, 2021. <https://doi.org/10.1145/3487552.3487827>

Execution time statistics

(a)

| Input Size | isWasmFaster | isJsFaster |
|------------|--------------|------------|
| XS | 40 | 1 |
| S | 39 | 2 |
| M | 23 | 18 |
| L | 25 | 16 |
| XL | 23 | 18 |

Average memory usages (in KB)


(b)

| Input Size | WASM | JS |
|------------|------------|--------|
| XS | 2,001.54 | 879.41 |
| S | 2,077.27 | 878.73 |
| M | 2,985.78 | 880.54 |
| L | 26,991.05 | 883.10 |
| XL | 100,943.88 | 889.20 |

- A key goal of WebAssembly to achieve near-native performance
- WebAssembly spec reports WebAssembly runs on average **10% slower** than native code

Group, WebAssembly Community, and Andreas Rossberg. “WebAssembly Specification,” n.d., 185.

WebAssembly vs. Native Code



Jangda, Abhinav, Bobby Powers, Emery D Berger, and Arjun Guha. "Not So Fast: Analyzing the Performance of WebAssembly vs. Native Code," n.d., 15.

Root cause analysis

```
1 void matmul (int C[NI][NJ],  
2             int A[NI][NK],  
3             int B[NK][NJ]) {  
4     for (int i = 0; i < NI; i++) {  
5         for (int k = 0; k < NK; k++) {  
6             for (int j = 0; k < NJ; j++) {  
7                 C[i][j] += A[i][k] * B[k][j];  
8             }  
9         }  
10    }  
11 }
```

Root cause analysis

```

1 xor r8d, r8d           #i <- 0
2 L1:                    #start first loop
3 mov r10, rdx
4 xor r9d, r9d           #k <- 0
5 L2:                    #start second loop
6 imul rax, 4*NK, r8
7 add rax, rsi
8 lea r11, [rax + r9*4]
9 mov rcx, -NJ           #j <- -NJ
10 L3:                   #start third loop
11 mov eax, [r11]
12 mov ebx, [r10 + rcx*4 + 4400]
13 imul ebx, eax
14 add [rdi + rcx*4 + 4*NJ], ebx
15 add rcx, 1            #j <- j + 1
16 jne L3                #end third loop
17
18 add r9, 1              #k <- k + 1
19 add r10, 4*NK
20 cmp r9, NK
21 jne L2                #end second loop
22
23 add r8, 1              #i <- i + 1
24 add rdi, 4*NJ
25 cmp r8, NI
26 jne L1                #end first loop
27 pop rbx
28 ret

```

```

1 mov [rbp-0x28], rax
2 mov [rbp-0x20], rdx
3 mov [rbp-0x18], rcx
4 xor edi, edi           #i <- 0
5 jmp L1'                #start first loop
6 L1:
7 mov ecx, [rbp-0x18]
8 mov edx, [rbp-0x20]
9 mov eax, [rbp-0x28]
10 L1':
11 imul r8d, edi, 0x1130
12 add r8d, eax
13 imul r9d, edi, 0x12c0
14 add r9d, edx
15 xor r11d, r11d         #k <- 0
16 jmp L2'                #start second loop
17 L2:
18 mov ecx, [rbp-0x18]
19 L2':
20 imul r12d, r11d, 0x1130
21 lea r14d, [r9+r11*4]
22 add r12d, ecx
23 xor esi, esi           #j <- 0
24 mov r15d, esi
25 jmp L3'                #start third loop
26 L3:
27 mov r15d, eax
28 L3':
29 lea eax, [r15+0x1]     #j <- j + 1
30 lea edx, [r8+r15*4]
31 lea r15d, [r12+r15*4]
32 mov esi, [rbx+r14*1]
33 mov r15d, [rbx+r15*1]
34 imul r15d, esi
35 mov ecx, [rbx+rdx*1]
36 add ecx, r15d
37 mov [rbx+rdx*1], ecx
38 cmp eax, NJ           #j < NJ
39 jnz L3                 #end third loop
40 add r11, 0x1           #k++
41 cmp r11d, NK          #k < NK
42 jnz L2                 #end second loop
43 add edi, 0x1          #i++
44 cmp edi, NI           #i < NI
45 jnz L1                 #end first loop
46 ret

```

- Native code uses 10 registers while WebAssembly uses 13 registers
- WebAssembly reserves r13 to point to an array of GC roots at all times
- WebASsembly uses r10 and xmm13 as dedicated scratch registers

- Extra jumps to avoid memory loads
- Stack overflow checks per function call
- Function table indexing checks

- Register issues can be ameliorated by improved implementations
- Stack overflow checks, indirect call checks, and reserved registers checks are necessary for WebAssembly's safety guarantees

Performance comparison between C and Rust compiled to WebAssembly

- 1) Matrix multiplication: 100X100 arrays with values 1-10,000
- 2) Insertion sort: int values between 0-999 in an array of size 1000
- 3) Add: A function that took two values and added them together


Medin, Magnus, and Ola Ringdahl. "Bachelor's Degree in Computer Science," n.d.,16.

Number of instructions for each test:

| Algorithm | get | set | add | store | load | offset | memory | tee |
|-----------|-----|-----|-----|-------|------|--------|--------|-----|
| Rust MM | 59 | 34 | 40 | 3 | 2 | 0 | 0 | 0 |
| C MM | 270 | 176 | 23 | 31 | 43 | 69 | 0 | 0 |
| Rust Sort | 21 | 11 | 12 | 3 | 2 | 0 | 0 | 8 |
| C Sort | 140 | 105 | 12 | 13 | 22 | 23 | 0 | 0 |
| Rust ADD | 2 | 0 | 3 | 0 | 0 | 0 | 3 | 0 |
| C ADD | 12 | 6 | 3 | 0 | 0 | 4 | 0 | 0 |


Languages compiled to WebAssembly

Matrix multiplication: Red bar - fetch time, Blue bar- execution time




Languages compiled to WebAssembly

Insertion sort: Red bar - fetch time, Blue bar- execution time



Languages compiled to WebAssembly

ADD: Red bar - fetch time, Blue bar- execution time



Takeaways:

- 1) WebAssembly performs differently when it compiles from different languages
- 2) Rust is preferable for more demanding calculations
- 3) For shorter and simpler tasks it is not as clear which language to choose
- 4) Performance of WebAssembly depends on the execution in browser engine rather than file transfer from the server